

---

# Modeling Cooperation in Heterogeneous Multi-Agent Systems

---

## Abstract

Modeling the complexity of global economics and politics requires capturing the diverse interactions of real-world actors with varying capabilities. To address this, we introduce a framework for running Multi-Agent Reinforcement Learning scenarios with heterogeneous agents. These agents vary in learning rates, risk tolerances, and observational abilities, allowing them to more accurately represent the diverse behaviors of real-world actors. Preliminary results suggest that sustainable cooperation can emerge even in the presence of significant asymmetries in agent capabilities. This highlights the potential of advanced computational models to help inform policies that improve global economic and political collaboration.

## 1 Introduction

Cooperation has been the key factor allowing humans to flourish Hare and Woods [2020], Bar-Yam [2000]. The complexity of modern societies often surpasses our evolutionary capacity for effective coordination, leading to social dilemmas where individual incentives and collective welfare are misaligned. These dilemmas manifest themselves in various domains, from the struggle to implement effective climate policies Kersting [2017], to the difficulties in agreeing to resolutions in order to deal with pandemics in a timely manner Glennerster et al. [2024]. Coordination problems are also prevalent in politics, where divergent interests and asymmetric information cause a misalignment of priorities among stakeholders Frieden and Lake [2015], leading to policy paralysis or, in extreme cases war Dafoe et al. [2020]. Similarly, in economics cooperation is often undermined by institutional failures Acemoglu and Robinson [2012]. Institutions may fail either because some members exploit others or because a lack of trust prevents cooperation, despite mutual cooperation maximizing collective welfare.

Recent advancements in economic modeling have increasingly focused on utilizing Markov Decision Processes (MDPs) to address these coordination problems Glennerster et al. [2024], Charpentier et al. [2021]. MDPs, coupled with Reinforcement Learning (RL) techniques like Temporal Difference (TD) learning, offer a promising framework for capturing the dynamic interactions among agents. In particular, these approaches allow for modeling strategic behavior where agents act based on current states without full knowledge of future outcomes. Agents in these models can adjust their strategies in response to the evolving economic landscape and other agents, much like governments, firms, and consumers in the global economy.

In this work, we introduce the novel capability of running TD simulations with heterogeneous agents. This method allows for the exploration of scenarios where heterogeneity—whether in terms of access to information, risk tolerance, or market reputation—plays a pivotal role in shaping outcomes. By simulating environments that reflect these disparities, we can better understand how different agent characteristics influence overall market dynamics and the potential for cooperation.

To model these different interaction scenarios, we utilize general-sum game environments, where the total payoff for all participants does not necessarily sum to zero. In contrast to zero-sum games, where one agent’s gain is exactly another agent’s loss, general-sum games allow for a range of outcomes

in which cooperation or competition can emerge. In these settings, cooperation may arise even in situations where agents have incentives to act in their own self-interest Barfuss [2022]. Over time, self-interested strategies can converge toward a common equilibrium that benefits both sides, even in the absence of a clear immediate incentive to cooperate, as is characteristic of general-sum games.

This paper is still a work-in-progress, and our current focus has primarily been on adjusting the agents' varying capacities to observe and react to their environment. Our early findings suggest that cooperation between heterogeneous agents that have a knowledge imbalance is feasible. This preliminary result indicates that even in games with heterogeneous agents, there is potential for cooperation, and it points to the broader applicability of these techniques in shaping more equitable economic outcomes.

## 2 Methods

### 2.1 Learning Algorithm

We employ the deterministic Expected SARSA Sutton and Barto [2018] algorithm to simulate the interactions of heterogeneous agents within a general-sum game scenario. Here we build upon previous work by Barfuss and Meylahn [2023].

Our framework considers two agents,  $i \in \{1, 2\}$ , where each agent selects from a set of possible actions,  $a^i \in \{c, d\}$ , representing a cooperating or defecting action, respectively. The joint action vector is denoted as  $\mathbf{a} = \{a^1, a^2\}$ . In the traditional SARSA algorithm the state  $s \in S$  is shared by both agents and together with the agent's actions, determines the reward  $r^i(s, \mathbf{a})$  received by agent  $i$  and the transition to the next state  $s' \in S$ .

However in our project, we then model the differences in state observability among agents by introducing an observed state  $\tilde{s} \in S$ . This clarifies the fact that one agent might not be observing the same state as another agent, even if they co-exist in the same environment.

At each time step  $t$ , agent  $i$  selects action  $a$  with a probability  $x_t^i(a|\tilde{s})$ , a frequency distribution over actions dependent on the observed state  $\tilde{s}$ . The action selection is guided by an  $\epsilon$ -greedy exploration strategy, where the agent chooses the action with the highest state-action value with probability  $1 - \epsilon_i$ . Next, the agent selects an action at random with probability  $\epsilon_i$  to encourage exploration. Specifically, for the two-action case, the action probabilities are defined as:

$$x_t^i(c|\tilde{s}) = \begin{cases} 1 - \frac{\epsilon_i}{2} & \text{if } q_t^i(c, a) > q_t^i(d, a), \\ \frac{\epsilon_i}{2} & \text{otherwise,} \end{cases} \quad (1)$$

where  $q_t^i(a, \tilde{s})$  represents the state-action value function for agent  $i$ , updated according to the temporal difference learning rule.

The state-action values are updated after each time step as follows:

$$q_{t+1}^i(\tilde{s}_t, a_t) = (1 - \alpha_i)q_t^i(\tilde{s}_t, a_t) + \alpha \left[ r_t^i + \delta \sum_a x_t^i(a|\tilde{s}_{t+1})q_t^i(\tilde{s}_{t+1}, a) \right], \quad (2)$$

where  $\alpha_i \in [0, 1]$  is the learning rate and  $\delta_i \in [0, 1]$  is the discount factor for agent  $i$ , which controls how much weight the agent places on future rewards. These parameters can be set differently for each agent to reflect individual differences in learning rates and future reward valuations.

Furthermore, the varying  $\epsilon_i$  and  $\delta_i$  parameters allow us to simulate agents with differing boldness in their decision-making processes, creating a parallel with the risk tolerance concept observed in real-world markets.

### 2.2 Environment

The environment we utilize to test our method is based on the Iterated Prisoner's Dilemma, a well-known framework for examining cooperation. This model serves as a representative example of how individual incentives can conflict with collective welfare, as it captures the tension between

cooperating for mutual benefit and defecting for individual gain. The reward matrix used for the two-agent case is given specified in the Appendix in Table 1.

For the model simulations, we utilize the framework of Multi-Agent Reinforcement Learning (MARL). While powerful, agent based models and Deep Reinforcement Learning (DRL) systems have significant disadvantages for improving the understanding of the learning agent’s emergent collective behavior: they are noisy, sometimes hard to explain, sample-inefficient, and computationally intense. We, therefore, employ an approach to MARL inspired by complex system science and evolutionary game theory: collective reinforcement learning dynamics (CRLD) Barfuss [2023].

CRLD is characterized by two types of idealization. First, CRLD uses dynamic learning equations as a computationally efficient alternative to traditional RL updates, allowing for simulations to be run on personal computers without the need for extensive computational resources. Second, CRLD can be used to study the principles behind emergent collective behavior in idealized, low-dimensional environments.

### 3 Results

For our initial experiments, we chose to use the IPD with a two-memory period as a test bed. Our approach leverages the concept of direct reciprocity in cooperation, as theorized by Schmid et al. [2021], in which agents base their cooperative decisions on personal historical interactions.

In a first experiment with 5,000 iterations we reliably found that agents that are more knowledgeable receives more rewards than less knowledgeable agents. See appendix for more details.




Figure 1: Distribution of average rewards for non self-aware agents. Agent 1 (blue) has complete observability while Agent 2 (yellow) has partial observability.




Figure 2: Distribution of average rewards for self-aware agents. Agent 1 (blue) has complete observability while Agent 2 (yellow) has partial observability.

In our second experimental setup involving 8,000 iterations with two agents. Agent 1 is always able to completely observe its environment, it has perfect memory and perfect understanding. Agent 2’s observational capabilities varied across two conditions: in the first, Agent 2 could observe the full market in the first game but in the second it loses the ability to observe its own actions (Table 2). In the second condition, Agent 2 retained knowledge of its own actions but could not observe Agent 1’s states (Table 3).

Our results, reveal striking differences in cooperative outcomes based on these observational asymmetries. Figure 1 shows the distribution of average rewards for the "Non Self-Aware Agent" scenario. We observe a strong tendency toward non-cooperative outcomes, with most interactions yielding low or zero rewards, showing that when an agent lack awareness of its own actions the results for cooperation are catastrophic. Figure 2 illustrates the outcomes for the case in which Agent 2 is a self-aware agent. This graph displays a broader distribution of rewards, with significant occurrences at higher values, suggesting that self-awareness facilitates better strategic decision-making and increases cooperation.

In a third experiment, an agent’s perception of the environment influences their roles as either exploiters or the exploited. When Agent 2 is only aware of instances when Agent 1 cooperates (Table 4), it tends to overestimate its own appropriate level of cooperation. This miscalculation leads to

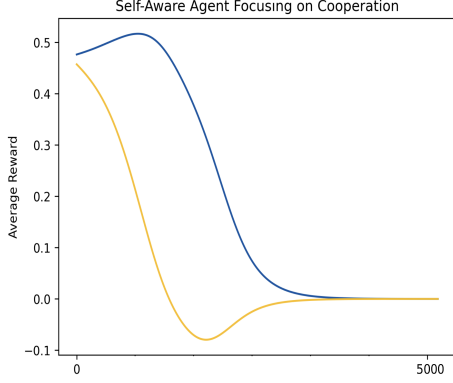


Figure 3: Average reward trends with Agent 2 (yellow) with partial observability focusing on cooperation and being exploited by Agent 1 (blue) with full observability.

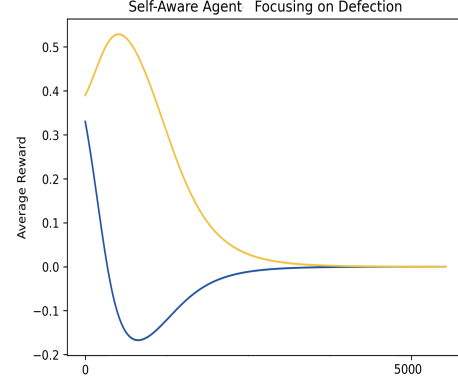


Figure 4: Average reward trends showing Agent 2 (yellow) with partial observability focusing on defection and exploiting Agent 1's (blue) despite Agent 1's greater awareness.

Agent 2 being exploited by Agent 1, as illustrated in Figure 3. Conversely, when Agent 2 is only aware of Agent 1 defections (Table 5), it becomes overly skeptical, choosing to exploit Agent 1.

## 4 Future work

Moving forward, we aim to further test our framework in a variety of strategic interaction settings, such as public goods games, cooperative bargaining and common pool resource games. We also hope to further analyze the strategies developed by these agents to gain deeper insights in the development of cooperation through the simple, low-dimensional scenarios typical of RL. Building on this understanding, we plan to integrate DRL algorithms into our stack to enhance our model's capability to handle more complex environments and possibly incorporate Large Language Model agents. We hope that this could make our models more realistic in order to aid policymakers in crafting more informed and effective policies. Our ultimate goal is to develop a comprehensive suite of tools that combine MDP and game-theoretical approaches to better support decision-making processes in economics and public policy.

## 5 Limitations

The IPD framework, by design, encapsulates a limited set of possible interactions and may not capture the full spectrum of strategic behaviors that agents exhibit in more nuanced scenarios. Moreover, the computational model assumes that all agents rationalize and make decisions based solely on predefined rules and their perception of the environment, which might not accurately reflect the often irrational and unpredictable nature of human decision-making. The model's outcomes are also highly dependent on the parameters chosen for the simulations, such as learning rates and discount factors, which might limit the generalizability of the findings to real-world applications, where such parameters are not as easily quantifiable or consistent.

## 6 Acknowledgments

This project was supported by Impact Academy and the University of Bonn. The authors don't have any conflict of interest to disclose.

## References

- Daron Acemoglu and James A. Robinson. *Why nations fail: The origins of power, prosperity, and poverty*. Currency, New York, 2012.
- Yaneer Bar-Yam. Complexity rising: From human beings to human civilization, a complexity profile. 2000.
- Wolfram Barfuss. Dynamical systems as a level of cognitive analysis of multi-agent learning. *Neural Computing and Applications*, 34(3):1653–1671, 2022. ISSN 1433-3058. doi: 10.1007/s00521-021-06117-0. URL <https://doi.org/10.1007/s00521-021-06117-0>.
- Wolfram Barfuss. pyCRLD: Collective reinforcement learning dynamics in python, 2023. URL <https://wbarfuss.github.io/pyCRLD/>.
- Wolfram Barfuss and Janusz M. Meylahn. Intrinsic fluctuations of reinforcement learning promote cooperation. *Scientific Reports*, 13(1):1309, 2023.
- Arthur Charpentier, Romuald Elie, and Carl Remlinger. Reinforcement learning in economics and finance. *Computational Economics*, pages 1–38, 2021.
- Allan Dafoe, Edward Hughes, Yoram Bachrach, Tantum Collins, Kevin R McKee, Joel Z Leibo, Kate Larson, and Thore Graepel. Open problems in cooperative ai. *arXiv preprint arXiv:2012.08630*, 2020.
- Jeffrey A. Frieden and David A. Lake. *World Politics: Interests, Interactions, Institutions*. W. W. Norton & Company, New York, 2015.
- Rachel Glennerster et al. Quantifying the social value of a universal covid-19 vaccine and incentivizing its development. Working Paper 32059, National Bureau of Economic Research, 2024.
- Brian Hare and Vanessa Woods. Survival of the friendliest. *Scientific American*, Aug 1 2020.
- Jan Kersting. *Stability of cooperation in the international climate negotiations: An analysis using cooperative game theory*. KIT Scientific Publishing, 2017.
- Laura Schmid, Krishnendu Chatterjee, Christian Hilbe, and Martin A. Nowak. A unified framework of direct and indirect reciprocity. *Nature Human Behaviour*, 5(10):1292–1302, 2021.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement learning: An introduction*. MIT Press, 2018.

## A Appendix / supplemental material

### A.1 Software

Code will be made available upon publication. A link to the repository will be provided in the camera-ready version.

### A.2 Parameters

For all experiments a constant learning rate of 0.1 and discount factor of 0.9 were used. Though it's possible to change these values as we have noted, in the interest of space we only presented findings maintaining these constant.

### A.3 Figures

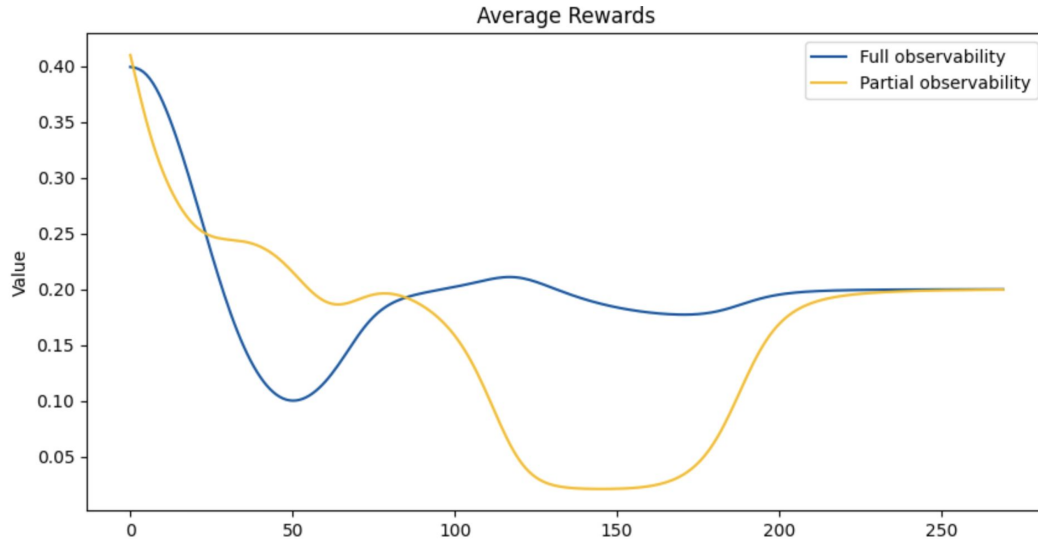


Figure 5: Average reward for agents within mixed observability environments, ran over 5000 iterations. We observe that more knowledgeable agents, receive better rewards on the long run than less knowledgeable agents. It is also clear from our analysis that exploitation is not a dominating behavior as agents tend towards mostly defecting or oscillating between cooperating and defecting.

### A.4 Tables

Table 1: Payoff Matrix for IPD

	c	d
c	(1, 1)	(1.2, -0.5)
d	(-0.5, 1.2)	(0, 0)

Table 2: 1st Experiment: States observed by agents in two iterations of the Prisoner's Dilemma

Agent 1, Game 1	Agent 2, Game 1	Agent 1, Game 2	Agent 2, Game 2
(C, C)	(C, C)	(C, C)	(C, *)
(C, C)	(C, D)	(C, C)	(C, *)
(C, C)	(D, C)	(C, C)	(C, *)
(C, C)	(D, D)	(C, C)	(C, *)
(C, D)	(C, C)	(C, D)	(C, *)
(C, D)	(C, D)	(C, D)	(C, *)
(C, D)	(D, C)	(C, D)	(C, *)
(C, D)	(D, D)	(C, D)	(C, *)
(D, C)	(C, C)	(D, C)	(D, *)
(D, C)	(C, D)	(D, C)	(D, *)
(D, C)	(D, C)	(D, C)	(D, *)
(D, C)	(D, D)	(D, C)	(D, *)
(D, D)	(C, C)	(D, D)	(D, *)
(D, D)	(C, D)	(D, D)	(D, *)
(D, D)	(D, C)	(D, D)	(D, *)
(D, D)	(D, D)	(D, D)	(D, *)

Table 3: 2nd Experiment: States observed by agents in two iterations of the Prisoner's Dilemma

Agent 1, Game 1	Agent 2, Game 1	Agent 1, Game 2	Agent 2, Game 2
(C, C)	(C, C)	(C, C)	(*, C)
(C, C)	(C, D)	(C, C)	(*, D)
(C, C)	(D, C)	(C, C)	(*, C)
(C, C)	(D, D)	(C, C)	(*, D)
(C, D)	(C, C)	(C, D)	(*, C)
(C, D)	(C, D)	(C, D)	(*, D)
(C, D)	(D, C)	(C, D)	(*, C)
(C, D)	(D, D)	(C, D)	(*, D)
(D, C)	(C, C)	(D, C)	(*, C)
(D, C)	(C, D)	(D, C)	(*, D)
(D, C)	(D, C)	(D, C)	(*, C)
(D, C)	(D, D)	(D, C)	(*, D)
(D, D)	(C, C)	(D, D)	(*, C)
(D, D)	(C, D)	(D, D)	(*, D)
(D, D)	(D, C)	(D, D)	(*, C)
(D, D)	(D, D)	(D, D)	(*, D)

Table 4: 3rd Experiment: States observed by agents in two iterations of the Prisoner's Dilemma

Agent 1, Game 1	Agent 2, Game 1	Agent 1, Game 2	Agent 2, Game 2
(C, C)	(C, C)	(C, C)	(C, C)
(C, C)	(C, D)	(C, C)	(C, D)
(C, C)	(D, C)	(C, C)	(*, C)
(C, C)	(D, D)	(C, C)	(*, D)
(C, D)	(C, C)	(C, D)	(C, C)
(C, D)	(C, D)	(C, D)	(C, D)
(C, D)	(D, C)	(C, D)	(*, C)
(C, D)	(D, D)	(C, D)	(*, D)
(D, C)	(C, C)	(D, C)	(C, C)
(D, C)	(C, D)	(D, C)	(C, D)
(D, C)	(D, C)	(D, C)	(*, C)
(D, C)	(D, D)	(D, C)	(*, D)
(D, D)	(C, C)	(D, D)	(C, C)
(D, D)	(C, D)	(D, D)	(C, D)
(D, D)	(D, C)	(D, D)	(*, C)
(D, D)	(D, D)	(D, D)	(*, D)

Table 5: 4th Experiment: States observed by agents in two iterations of the Prisoner's Dilemma

Agent 1, Game 1	Agent 2, Game 1	Agent 1, Game 2	Agent 2, Game 2
(C, C)	(C, C)	(C, C)	(*, C)
(C, C)	(C, D)	(C, C)	(*, D)
(C, C)	(D, C)	(C, C)	(D, C)
(C, C)	(D, D)	(C, C)	(D, D)
(C, D)	(C, C)	(C, D)	(*, C)
(C, D)	(C, D)	(C, D)	(*, D)
(C, D)	(D, C)	(C, D)	(D, C)
(C, D)	(D, D)	(C, D)	(D, D)
(D, C)	(C, C)	(D, C)	(*, C)
(D, C)	(C, D)	(D, C)	(*, D)
(D, C)	(D, C)	(D, C)	(D, C)
(D, C)	(D, D)	(D, C)	(D, D)
(D, D)	(C, C)	(D, D)	(*, C)
(D, D)	(C, D)	(D, D)	(*, D)
(D, D)	(D, C)	(D, D)	(D, C)
(D, D)	(D, D)	(D, D)	(D, D)



## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract informs what novelty the paper brings (heterogeneous agents within MARL contexts). It also emphasizes the intention to use the software as a tool to assist policymakers in decision-making, while clarifying that it is not intended to be used as a standalone solution. The introduction further elaborates on these key points.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: The paper includes a detailed "Limitations" section that discusses several key constraints and assumptions inherent to the study. It acknowledges the simplifications made by employing the IPD to model complex real-world interactions and the potential limitations of generalizing the findings due to specific model parameters like learning rates and discount factors. The discussion also addresses the assumptions regarding agent rationality and decision-making processes.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best

judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory Assumptions and Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [NA]

Justification: The focus of the paper is primarily on empirical evaluation of MARL scenarios using the IPD. While the paper discusses theoretical underpinnings and utilizes well-established algorithms, it does not introduce new theoretical results that require formal proofs.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper provides details about the experimental setups, including the parameters used (Table 1), learning rates, discount factors, the nature of the agents involved, and the simulation environments and software utilized. We provide Jupyter Notebooks displaying the learning rates and discount factors, as well as containing the code to reproduce all of the results found in the paper.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.

- (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The paper ensures open access to both the data and code necessary for reproducing the main experimental results. Jupyter Notebooks are provided, which contain the code used in the experiments. There is also extensive documentation on how to use the software we have used to run our tests (CRLD).

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental Setting/Details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: The paper doesn’t use training nor tests, however all of the important variables in order to replicate the results are both included in the Appendix and the Jupyter Notebooks.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment Statistical Significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Yes, the paper indicates the number of runs for each experiment where it's relevant to do so. Specifically we specify the number of runs for experiment 1 (Figure 1), experiment 2 (Figure 2 and Figure 3) and experiment 3 (Figure 4 and Figure 5).

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments Compute Resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We talk about how we decrease compute time by using the CRDT software, which allows for users to run computations in their personal computer.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code Of Ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research presented in this paper fully conforms to the NeurIPS Code of Ethics.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

#### 10. **Broader Impacts**

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA] .

Justification: This research presents a foundational framework, laying the groundwork for potential future applications. While the authors acknowledge that advanced iterations of this technology could have significant societal implications, both positive and negative, the current state of development is primarily theoretical and not yet applicable to real-world scenarios in a manner that could result in tangible societal impact.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

#### 11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA] .

Justification: Similar to the answer to the above question, this is a foundational theoretical work and in its current state doesn't offer misuse risks.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: Yes. The only licensed product we use is CRLD and it's under the GNU license.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

## 13. New Assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The package includes Jupyter Notebooks with tutorials that progress from simple to complex examples. It also features a website with comprehensive documentation on how to use the API. We hope this library will be adopted by others and expanded by the community.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

## 14. Crowdsourcing and Research with Human Subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[NA\]](#) .

Justification: No human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.

- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

**15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA] .

Justification: No human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.